

Wed Oct 21 2020

Nuru Nabiyeu 2642819

Matthias Debernardini 2622050

Det Shahu 2611771

Milestone 5 for Userspace TCP Stack

To understand how benchmarks are done we were tasked with doing one so that we can see the difficulties that come from trying to clock a computer system.

The first step involved choosing the platform to run the benchmark on. For that we decided to use a PC with the following specifications; AMD Ryzen 2700x, 16 GB of 3200MHz CL14 DDR4 RAM, 120GB SATAIII SSD running Ubuntu 18.04.5 LTS. All samples were taken from this machine consecutively. The second step involved choosing how to query the system clock. For this we picked the monotonic clock in the time library that ships with the linux kernel. We chose to use a wall clock rather than a user clock because we are interested in how the application uses the userspace library rather than how the CPU executes the code. The next step was deciding where to start the time and when to end it. We were interested in measuring the latency that occurs in the 3 way TCP handshake. To measure this, we placed the clocks in the client applications at the beginning and at the end of the exchange. Nothing was changed neither in the application nor in the userspace library. The difference between the clocks was printed to standard output, which is then redirected to a file with a timestamp of the sample. The code can be found in the appendix. The server and client were then ran multiple times.

Once we collected the data, we had to decide which samples were actually representative of the actual latency between the server and the client. It is unfortunately insufficient to just measure the time and then report one quantity. The environment that the code is running in is constantly changing, cache lines are invalidated at seemingly random intervals, the clock frequency is dynamically changing depending on the load (to increase energy efficiency and overall performance), pagefaults can occur (or not). This is too much information to try to distill into one number. Therefore, to really estimate the run-time of the exchange, we must also report the median of a population sample, the standard deviation and the uncertainty.

Because of how wildly variable the latency can be, outliers will need to be methodically dealt with. Our method for dealing with outliers involves calculating the Median-Absolute-Deviation, using the MAD to decide if a data point is an outlier. This step included deciding how many deviations away from the median we tolerate. In our case it was 3, this was chosen manually by increasing the amount of deviations and seeing when the changes in the number

of outliers was small relative to the number of samples. The uncertainty was then calculated by dividing the MAD by the square root of non outlier values

Statistic	Nanoseconds
Average Latency	51136
Median Latency	52665
95 Percentile	55601
99 Percentile	61627
Max Observed Latency	62051
Min Observed Latency	43060
Uncertainty	461
Sample Size	74

Uncertainty 461 Average 51136 Median 52665 Max Observed 62051 Min Observed 43060 95%% percentile: 55601 99%% percentile: 61627




Figure 1: CDF

```
/* measure monotonic time */
uint64_t delta;
struct timespec start, end;
clock_gettime(CLOCK_MONOTONIC, &start); /* mark start time */
```

```

/* send test buffer */
while (so_far < TEST_BUF_SIZE) {
    ret = send(server_fd, tx_buffer + so_far, TEST_BUF_SIZE - so_far, 0);
    if (0 > ret) {
        printf("Error: send failed with ret %d and errno %d \n", ret, errno);
        return -ret;
    }
    so_far += ret;
    printf("\t [send loop] %d bytes, looping again, so_far %d target %d \n", ret, so_far, TEST_BUF_SIZE);
}

printf("OK: buffer sent successfully \n");
printf("OK: waiting to receive data \n");
// receive test buffer
so_far = 0;
while (so_far < TEST_BUF_SIZE) {
    ret = recv(server_fd, rx_buffer + so_far, TEST_BUF_SIZE - so_far, 0);
    if (0 > ret) {
        printf("Error: recv failed with ret %d and errno %d \n", ret, errno);
        return -ret;
    }
    so_far += ret;
    printf("\t [receive loop] %d bytes, looping again, so_far %d target %d \n", ret, so_far, TEST_BUF_SIZE);
}
clock_gettime(CLOCK_MONOTONIC, &end);    /* mark the end time */

delta = BILLION * (end.tv_sec - start.tv_sec) + end.tv_nsec - start.tv_nsec;
printf("<<BENCHMARK>> %llu\n", (long long unsigned int) delta);

```

Reproducing the results

- Install just to expedite the benchmark otherwise copy and paste the commands found in the justfile directly in the shell
- just prep
- just benchmark
- cd plotter
- python3 main.py

http://blogs.perl.org/users/steffen_mueller/2010/09/your-benchmarks-suck.html

Goal

In order to get an estimate of the expectation value and its uncertainty, we need a model of the underlying distribution: Assuming these are comparatively rare and typically cause extraordinarily long run-times (as opposed to extraordinarily low run-times), we arrive at an overall model of having a central, smoothish normal distribution with a few outliers towards long run-times.